

## Table of Contents

<b>1. Background .....</b>	<b>2</b>
<b>2. Problem Statement.....</b>	<b>2</b>
<b>3. Exploratory Data Analysis .....</b>	<b>3</b>
3.1 EDA on Raw data .....	3
3.2 EDA on preprocessed data .....	6
<b>4. Data Pre-processing .....</b>	<b>12</b>
<b>5. Detailed Data Visualization .....</b>	<b>18</b>
<b>6. Data Preparation.....</b>	<b>26</b>
<b>7. Statistical Analysis:.....</b>	<b>30</b>
<b>8. Machine Learning Models (Training and Evaluation).....</b>	<b>32</b>
8.1 Linear Regression Model.....	32
8.2 Random Forest.....	32
8.3 Support Vector Machine (SVM) .....	33
8.4 Gradient Boosting Machine.....	33
<b>9. Model Comparison and Selection .....</b>	<b>34</b>
<b>10. Adequate interpretation and justification of results .....</b>	<b>36</b>
<b>11. Practical Applications for Organizational Decision-Making .....</b>	<b>37</b>
<b>12. Conclusion.....</b>	<b>37</b>
<b>13. References .....</b>	<b>39</b>

## **1. Background**

The quick progress in weather science has boosted our skill to predict weather and grasp environmental states. Yet, it's still tricky to forecast local weather factors like heat, air pressure, and ground wetness. These elements matter a lot for fields such as farming, handling disasters, and public health where spot-on weather forecasts can guide smarter choices.

The Weather Research and Forecasting (WRF) model is a popular tool to simulate weather at different levels, from worldwide to smaller areas. Our project zeros in on Scotland, a place known for its changing climate due to where it sits and its land shape. Getting weather forecasts right in this region is key to get ready for harsh weather and to manage nature's resources well.

## **2. Problem Statement**

Weather factors have a huge impact on our everyday lives and the economy. This project aims to boost how well we can predict key weather and environmental conditions in Scotland using the Weather Research and Forecast(WRF) data.

The primary objective includes:

### **1. Predicting Weather Variables:**

To develop a machine learning models to predict critical weather variables such as surface temperature(TSK), surface pressure(PSFC) and soil moisture(SMOIS) based on the other meteorological inputs like wind speed(U10, V10), humidity(Q2), precipitation(RAINC, RAINCC) and other relevant variable.

### **2. Daytime vs Nighttime analysis:**

To perform statistical analysis, such as T-tests to evaluate the different meteorological variables, such as surface pressure(PSFC) between daytime and nighttime, in order to understand daily variation.

This issue matters because better weather forecasts can help different areas in Scotland. Take farmers, for instance. If they know how warm it'll be and how wet the soil is, they can water their crops just right and plan what to grow. Also, if we can predict air pressure well, we can warn people about bad weather sooner. This project uses smart computer programs to make weather predictions more accurate. It also wants to teach us more about how the weather works in Scotland.

The main aim is to create a weather prediction tool people can count on. Weather experts, scientists, and government officials could use it to make smart choices. This could help Scotland deal with changing weather patterns and keep people safer overall.

### 3. Exploratory Data Analysis

EDA plays a key role as the first step in any data analysis project. It helps us grasp the underlying patterns, relationships, and distributions within the dataset. In this project, EDA had an impact on summarizing the main statistical properties of the data, showing distributions of weather variables, and spotting links between them. For example, we looked into how factors like temperature at skin level (TSK), surface pressure (PSFC), and soil moisture (SMOIS) interact with other weather elements such as wind speed (U10, V10), humidity (Q2), and rainfall (RAIN, RAINNC). EDA proved useful to find possible outliers, check normality, and get a handle on the data's structure. This information guided our choice of modeling approach later on.

#### 3.1 EDA on Raw data

##### Loading necessary libraries

```
library(dplyr)
library(readr)
library(ggplot2)
library(corrplot)
library(here)
library(stringr)
```

##### Loading the dataset

```
data <- read_csv("C:/Users/Ahmad Afzal/Desktop/WRFdata_May2023.csv", col_names = FALSE)
```

##### Assigning second row as a header and removing the second row

Assigning appropriate column names is essential for better data analysis and processing. Removing the row used for column names prevents duplication.

```
# Step 2: Assign the second row as the header
colnames(data) <- as.character(data[2, ])

# Step 3: Remove the second row from the data (to avoid duplication in the rows)
data <- data[-2, ]
```

##### Overview of the dataset

Initial exploration of the dataset helps in understanding the data types, structure, and distribution of variables. Summary statistics provide a quick overview of the central tendency and dispersion of the data.

```
print(dim(data)) # Dimensions of the dataset
print(str(data)) # Structure of the dataset
print(summary(data)) # Summary statistics for each column
```

```
> print(dim(data)) # Dimensions of the dataset
[1] 5453 2482
```

**Figure 1: Dimensions of Initial data**

```
> print(summary(data)) # Summary statistics for each column
```

X1	X2	X3	X4	X5	X6
X7	X8	X9	X10	X11	X12
X13	X14	X15	X16	X17	X18
X19	X20	X21	X22	X23	X24
X25	X26	X27	X28	X29	X30
X31	X32	X33	X34	X35	X36
X37	X38	X39	X40	X41	X42
X43	X44	X45	X46	X47	X48
X49	X50	X51	X52	X53	X54
X55	X56	X57	X58	X59	X60
X61	X62	X63	X64	X65	X66
X67	X68	X69	X70	X71	X72
X73	X74	X75	X76	X77	X78
X79	X80	X81	X82	X83	X84
X85	X86	X87	X88	X89	X90
X91	X92	X93	X94	X95	X96
X97	X98	X99	X100	X101	X102
X103	X104	X105	X106	X107	X108
X109	X110	X111	X112	X113	X114
X115	X116	X117	X118	X119	X120
X121	X122	X123	X124	X125	X126
X127	X128	X129	X130	X131	X132
X133	X134	X135	X136	X137	X138
X139	X140	X141	X142	X143	X144
X145	X146	X147	X148	X149	X150
X151	X152	X153	X154	X155	X156
X157	X158	X159	X160	X161	X162
X163	X164	X165	X166	X167	X168
X169	X170	X171	X172	X173	X174
X175	X176	X177	X178	X179	X180

**Figure 2: Initial Summary of Data**

## Missing data analysis

Analyzing missing data is crucial for determining whether the data is missing at random or systematically. Visualization helps in quickly identifying variables with significant missing data.

```
missing_data_summary <- colSums(is.na(data))
print(missing_data_summary)
```

ta.

```
> # Missing data analysis
> missing_data_summary <- colSums(is.na(data))
> print(missing_data_summary)
```

X1	X2	X3	X4	X5	X6	X7	X8	X9
156	162	150	168	154	153	156	157	173
X10	X11	X12	X13	X14	X15	X16	X17	X18
124	148	158	160	148	162	153	148	146
X19	X20	X21	X22	X23	X24	X25	X26	X27
176	152	163	152	154	171	165	137	154
X28	X29	X30	X31	X32	X33	X34	X35	X36
169	178	166	161	154	150	161	160	148
X37	X38	X39	X40	X41	X42	X43	X44	X45
168	173	167	177	161	166	170	151	160
X46	X47	X48	X49	X50	X51	X52	X53	X54
140	176	171	154	184	172	153	158	153
X55	X56	X57	X58	X59	X60	X61	X62	X63
168	168	145	167	161	171	182	158	147
X64	X65	X66	X67	X68	X69	X70	X71	X72
157	143	163	135	170	150	164	174	169
X73	X74	X75	X76	X77	X78	X79	X80	X81
161	163	160	180	171	168	159	159	167
X82	X83	X84	X85	X86	X87	X88	X89	X90
185	160	159	166	163	142	180	184	147
X91	X92	X93	X94	X95	X96	X97	X98	X99

**Figure 3: Missing data summary**

```
missing_data_plot <- data.frame(Variable = names(missing_data_summary),
                                Missing_Count = missing_data_summary) %>%
  ggplot(aes(x = reorder(Variable, -Missing_Count), y = Missing_Count)) +
  geom_bar(stat = "identity", fill = "tomato") +
  coord_flip() +
  labs(title = "Missing Data Summary", x = "Variables", y = "Count of Missing Values") +
  theme_minimal()

print(missing_data_plot)
```




Figure 4: Bar plot to visualize the distribution of missing values

```
print(head(data))
```

```
> print(head(data))
# A tibble: 6 × 2,482
  XLAT  XLONG  TSK    PSFC  ` "U10" `  ` "V10" `  ` "Q2" `  RAINC
  <chr>  <chr>  <chr>  <chr>  <chr>      <chr>      <chr>      <chr>
1 X      X.1    X01.05... X.2    X.3        X.4        X.5        X.6
2 48.871 -11.221 NA      1014... 6.9        5.5        0.006... 0.0
3 49.010 -11.240 285.2    1013... 7.0        5.8        0.006... 0.0
4 49.149 -11.259 285.2    1013... 7.0        6.0        0.006... 0.0
5 49.288 -11.278 285.2    1013... 7.0        6.3        0.006... 0.0
6 49.427 -11.298 285.2    1012... 7.0        6.6        0.006... 0.0
# i 2,474 more variables: RAINNC <chr>, SNOW <chr>,
```

Figure 5: 1<sup>st</sup> five rows of data

## 3.2 EDA on preprocessed data

```
print(str(newdata_copy))
print(head(newdata_copy))
print(summary(newdata_copy))
```

```
> print(head(newdata_copy))
# A tibble: 6 × 11
  TSK    PSFC X.U10. X.V10.    X.Q2. RAINC RAINNC  SNOW  TSLB SMOIS
  <dbl> <dbl> <dbl> <dbl>    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1  273.  97392    3.1    0.6 0.00393    NA     0     0  278.  0.270
2  272.  96298    2.6    1.5 0.00362     0     0     0  278.  0.246
3  274.  99340   -0.9    3.6 0.00389     0     0     0  278.  0.264
4  282. 100477    1.2    7.7 0.00524     0     0     0  273.  1
5  272.  98697    0.6    0.7 0.00367     0     0     0  280.  0.287
6  274.  99203   -0.3    3.5 0.00422     0     0     0  278.  0.283
# 1 more variable: DATETIME <dtm>
~ |
```

Figure 6: 1<sup>st</sup> five rows of preprocessed data

```
> print(str(newdata_copy))
tibble [2,996 × 11] (S3: tbl_df/tbl/data.frame)
 $ TSK      : num [1:2996] 273 272 274 282 272 ...
 $ PSFC      : num [1:2996] 97392 96298 99340 100477 98697 ...
 $ X.U10.    : num [1:2996] 3.1 2.6 -0.9 1.2 0.6 -0.3 -0.7 0.9 0 1.3 ...
 $ X.V10.    : num [1:2996] 0.6 1.5 3.6 7.7 0.7 3.5 3.3 3 3.9 2.6 ...
 $ X.Q2.     : num [1:2996] 0.00393 0.00362 0.00389 0.00524 0.00367 0.00422 0.00408 0.00412 0.00465 0.00382 ...
 $ RAINC     : num [1:2996] NA 0 0 0 0 0 0 0 0 ...
 $ RAINNC    : num [1:2996] 0 0 0 0 0 0 0 0 0 ...
 $ SNOW      : num [1:2996] 0 0 0 0 0 0 0 0 0 ...
 $ TSLB      : num [1:2996] 278 278 278 273 280 ...
 $ SMOIS     : num [1:2996] 0.27 0.246 0.264 1 0.287 ...
 $ DATETIME : POSIXct[1:2996], format: "2018-05-01 00:00:00" "2018-05-01 00:00:00" ...
 NULL
~ |
```

Figure 7: Preprocessed data

```
> print(summary(newdata_copy))

      TSK      PSFC      X.U10.      X.V10.
Min.   :271.6   Min.   : 92974   Min.   : -8.100   Min.   : -10.600
1st Qu.:279.7   1st Qu.: 97671   1st Qu.:  0.300   1st Qu.:  2.400
Median :281.5   Median : 99095   Median :  2.500   Median :  4.200
Mean    :281.9   Mean    : 98937   Mean    :  2.315   Mean    :  4.141
3rd Qu.:283.4   3rd Qu.:100362   3rd Qu.:  4.500   3rd Qu.:  6.100
Max.    :298.0   Max.    :102423   Max.    :11.000   Max.    :15.300

      X.Q2.      RAINC      RAINNC      SNOW
Min.   :0.00309   Min.   :0.0000   Min.   : 0.0    Min.   :0.00000
1st Qu.:0.00508   1st Qu.:0.0000   1st Qu.: 0.0    1st Qu.:0.00000
Median :0.00610   Median :0.0000   Median : 0.2    Median :0.00000
Mean    :0.00598   Mean    :0.1084   Mean    : 1.3    Mean    :0.01046
3rd Qu.:0.00690   3rd Qu.:0.0000   3rd Qu.: 1.4    3rd Qu.:0.00000
Max.    :0.00929   Max.    :5.3000   Max.    :16.7   Max.    :1.40000
      NA's :1

      TSLB      SMOIS      DATETIME
Min.   :273.2   Min.   :0.2220   Min.   :2018-05-01 00:00:00.00
1st Qu.:273.2   1st Qu.:0.2609   1st Qu.:2018-05-03 12:00:00.00
Median :279.6   Median :0.2798   Median :2018-05-06 00:00:00.00
Mean    :279.2   Mean    :0.4666   Mean    :2018-05-06 00:02:24.19
3rd Qu.:282.4   3rd Qu.:1.0000   3rd Qu.:2018-05-08 12:00:00.00
Max.    :290.8   Max.    :1.0000   Max.    :2018-05-11 00:00:00.00

~ |
```

Figure 8: Summary of data

## Descriptive statistics for numerical values

```
print(summary(newdata_copy[, sapply(newdata_copy, is.numeric)]))
```

```
> print(summary(newdata_copy[, sapply(newdata_copy, is.numeric)]))
      TSK      PSFC      X.U10.      X.V10.
Min.   :271.6   Min.    : 92974   Min.    :-8.100   Min.    :-10.600
1st Qu.:279.7   1st Qu.: 97671   1st Qu.: 0.300   1st Qu.:  2.400
Median :281.5   Median : 99095   Median : 2.500   Median :  4.200
Mean   :281.9   Mean    : 98937   Mean    : 2.315   Mean    :  4.141
3rd Qu.:283.4   3rd Qu.:100362   3rd Qu.: 4.500   3rd Qu.:  6.100
Max.   :298.0   Max.    :102423   Max.    :11.000   Max.    : 15.300

      X.Q2.      RAINC      RAINNC      SNOW
Min.   :0.00309   Min.    :0.0000   Min.    : 0.0    Min.    :0.00000
1st Qu.:0.00508   1st Qu.:0.0000   1st Qu.: 0.0    1st Qu.:0.00000
Median :0.00610   Median :0.0000   Median : 0.2    Median :0.00000
Mean   :0.00598   Mean    :0.1084   Mean    : 1.3    Mean    :0.01046
3rd Qu.:0.00690   3rd Qu.:0.0000   3rd Qu.: 1.4    3rd Qu.:0.00000
Max.   :0.00929   Max.    :5.3000   Max.    :16.7    Max.    :1.40000
      NA's :1

      TSLB      SMOIS
Min.   :273.2   Min.    :0.2220
1st Qu.:273.2   1st Qu.:0.2609
Median :279.6   Median :0.2798
Mean   :279.2   Mean    :0.4666
3rd Qu.:282.4   3rd Qu.:1.0000
Max.   :290.8   Max.    :1.0000
```

Figure 9: Descriptive summary of numerical values

## Checking missing values

```
print(colMeans(is.na(newdata_copy)))
```

```
> print(colMeans(is.na(newdata_copy)))
      TSK      PSFC      X.U10.      X.V10.      X.Q2.
0.0000000000 0.0000000000 0.0000000000 0.0000000000 0.0000000000
      RAINC      RAINNC      SNOW      TSLB      SMOIS
0.0003337784 0.0000000000 0.0000000000 0.0000000000 0.0000000000
      DATETIME
0.0000000000
```

Figure 10: Missing values count

## Explore data distribution using Histogram

```
numeric_vars <- names(newdata_copy)[sapply(newdata_copy, is.numeric)]
for (var in numeric_vars) {
  print(
    ggplot(newdata_copy, aes_string(x = var)) +
    geom_histogram(bins = 30, color = "black", fill = "lightblue") +
    labs(title = paste("Distribution of", var), x = var, y = "Frequency") +
    theme_bw()
  )
}
```

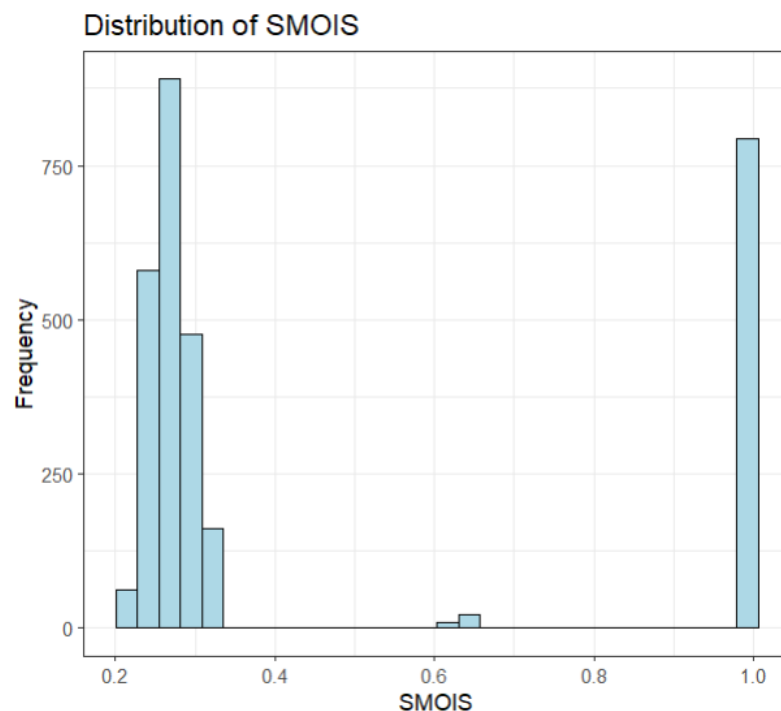


Figure 11: Data Distribution Histogram

## Identify and report potential outliers:

```
# Identify outliers with z-scores (more than 3 standard deviations from the mean)
outliers <- lapply(newdata_copy[numeric_vars], function(x) {
  abs_z_scores <- abs((x - mean(x, na.rm = TRUE)) / sd(x, na.rm = TRUE))
  return(which(abs_z_scores > 3))
})
```

```
# Report columns containing outliers based on z-scores
outlier_columns <- names(outliers)[sapply(outliers, length) > 0]
if (length(unlist(outliers)) > 0) {
  cat("Potential outliers identified in columns:", outlier_columns, "\n")
} else {
  cat("No potential outliers identified based on z-scores.\n")
}
```

```
Potential outliers identified in columns: TSK PSFC X.U10. X.V1
0. RAINC RAINNC SNOW
```

Figure 12: Columns containing outliers



## Daytime vs Nighttime Pressure(PSFC) Analysis

A t-test helps determine whether there is a statistically significant difference in surface pressure between daytime and nighttime, providing insights into diurnal patterns.

```
# Daytime vs. Nighttime Pressure (PSFC) Analysis
data_df <- data1_clean
data_df$Hour <- as.numeric(format(data_df$DATETIME, "%H"))
data_df$Day_Night <- ifelse(data_df$Hour >= 6 & data_df$Hour < 18, "Daytime", "Nighttime")

# Subset data and perform t-test
daytime_data <- filter(data_df, Day_Night == "Daytime")
nighttime_data <- filter(data_df, Day_Night == "Nighttime")
t_test_result <- t.test(daytime_data$PSFC, nighttime_data$PSFC)
print(t_test_result)

> print(t_test_result)

Welch Two Sample t-test

data: daytime_data$PSFC and nighttime_data$PSFC
t = -1.0254, df = 2993.9, p-value = 0.3053
alternative hypothesis: true difference in means is not equal
to 0
95 percent confidence interval:
 -200.19014  62.70907
sample estimates:
mean of x mean of y
 98902.91  98971.65
```

*Figure 13: t-test result*

## Correlation analysis between SMOIS and TSK

Correlation between soil moisture and skin temperature can reveal important insights into the land-atmosphere interaction and surface energy balance.

```
# Correlation Analysis between SMOIS and TSK
correlation <- cor(data_df$SMOIS, data_df$TSK, use = "complete.obs")
print(correlation)
corrplot(cor(data_df[, c("SMOIS", "TSK")], use = "complete.obs"), method = "circle")

> print(correlation)
[1] -0.05095017
> corrplot(cor(data_df[, c("SMOIS", "TSK")], use = "complete.o
bs"), method = "circle")
> |
```

*Figure 14: Correlation between SMOIS and TSK*




Figure 15: Correlation Plot between SMOIS and TSK

## Wind speed analysis

Analyzing wind speed helps in understanding the wind patterns, which are critical for weather prediction, especially in the context of renewable energy sources like wind power.


```
# Mean Wind Speed Analysis
data_df$Date <- as.Date(data_df$DATETIME)
mean_wind_speed <- aggregate(WIND_SPEED ~ Date, data = data_df, FUN = mean)
print(mean_wind_speed)
```

```
> print(mean_wind_speed)
      Date WIND_SPEED
1 2018-04-30   3.705616
2 2018-05-01   6.882601
3 2018-05-02   6.569122
4 2018-05-03   6.583378
5 2018-05-04   6.888547
6 2018-05-05   5.845709
7 2018-05-06   4.634054
8 2018-05-07   4.577939
9 2018-05-08   5.757703
10 2018-05-09   6.750845
11 2018-05-10   5.904093
```

Figure 16: Mean of Wind Speed

## Distribution of Wind Speed

```
# Plot the distribution of wind speed
ggplot(data_df, aes(x = WIND_SPEED)) +
  geom_histogram(binwidth = 1, fill = "skyblue", color = "black") +
  labs(title = "Distribution of Wind Speed", x = "WIND_SPEED (km/h)", y = "Frequency") +
  theme_minimal()
```



**Figure 17: Distribution of Wind Speed**

## 4. Data Pre-processing

Data preprocessing is a crucial step to prepare the raw dataset for machine learning models. It involved several tasks, including:

**Handling Missing Values:** Rows with missing data in critical columns (e.g., TSK, PSFC, SMOIS, WIND\_SPEED, and RAINC) were removed to ensure the integrity of the analysis. This step was vital to prevent biases and inaccuracies in model predictions.

**Feature Engineering:** Additional features were created, such as converting the DATETIME column to a proper date-time format and categorizing times of the day into Day and Night to facilitate diurnal analysis. We have just extracted the data of 11 days and specific to that of Scotland

**Data Splitting:** The cleaned dataset was split into training (70%) and testing (30%) sets to evaluate model performance accurately. This step ensures that the models are trained on a representative sample and validated on unseen data to prevent overfitting.

**Standardization:** Numeric features were standardized to ensure that all variables contribute equally to the model's learning process, improving the overall performance and convergence of certain algorithms.

These preprocessing steps ensured that the data was in the best possible shape for developing robust and accurate predictive models, forming the foundation for the subsequent machine learning tasks.

### Filtering out Bolton's coordinates

Geographical filtering ensures that the data is relevant to the area of interest, and sampling helps in maintaining a manageable dataset size for analysis.

```
# Define Bolton's coordinates (approximately)
bolton_lat <- 53.5789
bolton_long <- -2.4292

# Set the conditions for filtering
min_latitude <- 49 # Example lower bound for latitude
max_latitude <- 54 # Example upper bound for latitude
min_longitude <- -11 # Example lower bound for longitude
max_longitude <- -7 # Example upper bound for longitude

colnames(data) <- make.names(names(data), unique = TRUE)

# Filter out rows that match Bolton's coordinates and fall within your specified latitude and longitude range
filtered_df <- data %>%
  filter(!(XLAT == bolton_lat & XLONG == bolton_long)) %>% # Exclude Bolton's coordinates
  filter(XLAT >= min_latitude & XLONG <= max_latitude) %>% # Select appropriate latitude range
  filter(XLAT >= min_longitude & XLONG <= max_longitude) # Select appropriate longitude range

# Ensure that at least 300 rows are selected
filtered_df <- filtered_df %>%
  sample_n(min(n(), 300)) # Randomly sample at least 300 rows or fewer if less than 300 rows are available

# Display the first few rows of the filtered dataset
print(head(filtered_df))
```

```
> # Display the first few rows of the filtered dataset
> print(head(filtered_df))
# A tibble: 6 x 2,482
  XLAT  XLONG  TSK  PSFC  X.U10.  X.V10.  X.Q2.  RAINC  RAINNC  SNOW  TSLB  SMOIS  TSK.1  PSFC.1  X.U10..1  X.V10..1  X.Q2..1  RAINC.1
<chr> <chr> <chr> <chr> <chr> <chr> <chr> <chr> <chr> <chr> <chr> <chr> <chr> <chr> <chr> <chr> <chr> <chr>
1 50.200 -3.565 283.5 1011... 6.1 -4.3 0.00... 0.0 0.0 NA 273.2 1.00... 283.5 101098 6.1 -0.1 0.00536 0.0
2 50.272 -11.2... 284.9 1011... 6.9 8.0 0.00... 0.0 0.0 0.0 273.2 1.00... NA 100744 7.0 10.3 0.00653 0.0
3 49.352 -5.712 284.4 1014... 6.0 0.3 0.00... 0.0 0.0 0.0 273.2 1.00... 284.4 101325 6.3 3.1 0.00557 0.0
4 54.776 -4.239 282.2 1004... 1.6 NA 0.00... 0.0 0.0 0.0 273.2 1.00... 282.2 NA 2.6 3.9 0.00565 0.0
5 55.051 -4.481 271.3 98026 2.0 0.1 0.00... 0.0 0.0 0.0 278.3 0.30... 272.1 97852 0.0 2.6 0.00405 0.0
6 59.557 -4.534 282.1 1004... -0.7 6.4 0.00... 0.0 NA 0.0 NA 1.00... 282.1 100219 -0.1 NA 0.00498 NA
# i 2,464 more variables: RAINNC.1 <chr>, SNOW.1 <chr>, TSLB.1 <chr>, SMOIS.1 <chr>, TSK.2 <chr>, PSFC.2 <chr>,
```

Figure 18: 1<sup>st</sup> five rows of data

## Setting up the correct columns in the dataset and adding the date\_time column

calculating appropriate datetime values and combining the results into a comprehensive data frame. This structured approach ensures that the data is correctly aligned with time and ready for further analysis, such as forecasting or modeling.

```
df <- read_csv("C:/Users/Ahmad Afzal/Desktop/cleaned_file.csv")

# Assuming your dataset is named 'df'
# Define the start date and time
start_datetime <- as.POSIXct("2018-05-01 00:00:00", format="%Y-%m-%d %H:%M:%S")

# Initialize an empty dataframe to store the results
result <- data.frame()

# Define the base columns to keep
base_columns <- c("XLAT", "XLONG", "TSK", "PSFC", "X.U10.", "X.V10.",
                  "X.Q2.", "RAINC", "RAINNC", "SNOW", "TSLB", "SMOIS")

# Loop over each set of columns
#This data is for the 10 days of data
for (i in 0:80) {
  # Create a copy of the base columns
  temp_df <- df[, base_columns]

  if (i > 0) {
    # Modify the column names to include the suffix
    suffix <- paste0(".", i)
    columns_with_suffix <- paste0(base_columns[-c(1:2)], suffix)
    temp_df[, -c(1:2)] <- df[, columns_with_suffix]
  }

  # Calculate the current datetime for this set
  current_datetime <- start_datetime + (i * 3 * 3600)

  # Format the datetime as required
  temp_df$date_time <- format(current_datetime, "%d.%m.%Y.%H.%M")

  # Bind the current dataframe to the result
  result <- bind_rows(result, temp_df)
}

# Display the result
head(result)
```

```

> # Display the result
> head(result)
  XLAT  XLONG  TSK  PSFC X.U10. X.V10.  X.Q2. RAINC
1 56.390 -0.779  NA 100571  3.5   -6.9    NA    0
2 56.559 -3.005 272.0 99379  2.3    0.4 0.00376    0
3 56.831 -2.749 272.6 97392  3.1    0.6 0.00393    NA
4 57.382 -4.760 272.3 96298  2.6    1.5 0.00362    0
5 56.877 -12.279 283.2 99604  4.6   14.3 0.00670    0
6 58.470 -5.040 274.4 99340  -0.9    3.6 0.00389    0
  RAINNC SNOW  TSLB  SMOIS      date_time
1  0.0    0 273.2 1.0000 01.05.2018.00.00
2  0.0    0 278.6 0.2779 01.05.2018.00.00
3  0.0    0 277.5 0.2696 01.05.2018.00.00
4  0.0    0 278.1 0.2457 01.05.2018.00.00
5  0.3    0 273.2 1.0000 01.05.2018.00.00
6  0.0    0 277.7 0.2640 01.05.2018.00.00
> |

```

Figure 19: 1<sup>st</sup> five rows of data

## Extracting the data having the coordinates of Scotland

We are just keeping the associated with the coordinates of Scotland to perform further tasks and operations

```
# Just keeping the coordinates of Scotland
```

```
scotland_lat_min <- 55.8
```

```
scotland_lat_max <- 58.6
```

```
scotland_long_min <- -6.5
```

```
scotland_long_max <- -2.0
```

```
# Filter for Coordinates within Scotland
```

```
scotland_data <- df[df$XLAT >= scotland_lat_min &
  df$XLAT <= scotland_lat_max &
  df$XLONG >= scotland_long_min &
  df$XLONG <= scotland_long_max, ]
```

```
# Print the filtered coordinates (Optional)
```

```
print(scotland_data)
```

```
# A tibble: 2,997 × 13
```

```

  XLAT XLONG  TSK  PSFC X.U10. X.V10.  X.Q2. RAINC RAINNC SNOW  TSLB
<dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1  56.6 -3.00 272  99379  2.3    0.4 0.00376    0      0    0 279.
2  56.8 -2.75 273.  97392  3.1    0.6 0.00393    NA     0    0 278.
3  57.4 -4.76 272.  96298  2.6    1.5 0.00362    0      0    0 278.
4  58.5 -5.04 274.  99340 -0.9    3.6 0.00389    0      0    0 278.
5  58.5 -5.56 282. 100477  1.2    7.7 0.00524    0      0    0 273.
6  56.0 -5.23 272.  98697  0.6    0.7 0.00367    0      0    0 280.
7  57.2 -5.77 274.  99203 -0.3    3.5 0.00422    0      0    0 278.
8  58.2 -5.29 274.  99758 -0.7    3.3 0.00408    0      0    0 278.
9  56.4 -5.98 274.  98914  0.9    3    0.00412    0      0    0 278.
10 58.5 -2.96 281. 100515  0     3.9 0.00465    0     NA    0 273.
# & 2 987 more rows

```

Figure 20: data representation

## Removing XLAT and XLONG columns from the dataset

As now these both columns are of no use so we will remove these columns from our dataset in order to make our data cleaner and more manageable to perform further tasks

```
data_without_coordinates <- data %>%
  select(-XLAT, -XLONG)

# Save the updated data to a new CSV file
write_csv(data_without_coordinates, "C:/Users/Ahmad Afzal/Desktop/Scotland_WRFdata_Without_Coordinates.csv")

# Print the first few rows of the updated data to the console
print(head(data_without_coordinates))
```

```
> # Print the first few rows of the updated data to the console
> print(head(data_without_coordinates))
# A tibble: 6 × 11
   TSK    PSFC X.U10. X.V10.   X.Q2. RAINC RAINNC SNOW  TSLB SMOIS
  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1  272.  99379    2.3    0.4 0.00376    0      0      0  279.  0.278
2  273.  97392    3.1    0.6 0.00393    NA      0      0  278.  0.270
3  272.  96298    2.6    1.5 0.00362    0      0      0  278.  0.246
4  274.  99340   -0.9    3.6 0.00389    0      0      0  278.  0.264
5  282. 100477    1.2    7.7 0.00524    0      0      0  273.  1
6  272.  98697    0.6    0.7 0.00367    0      0      0  280.  0.287
# i 1 more variable: date_time <chr>
> |
```

Figure 21: 1<sup>st</sup> five rows without coordination

## Renaming date\_time column as DATETIME

```
# Check if 'date_time' column exists and rename it
if ("date_time" %in% colnames(new_data)) {
  colnames(new_data)[colnames(new_data) == "date_time"] <- "DATETIME"
} else {
  stop("The 'date_time' column is missing in the dataset.")
}

# Print the first few rows of the updated new_data to check changes
print(head(new_data))
```

```
> # Print the first few rows of the updated new_data to check changes
> print(head(new_data))
# A tibble: 6 × 11
   TSK    PSFC X.U10. X.V10.    X.Q2. RAINC RAINNC  SNOW  TSLB SMOIS DATETIME
<dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <chr>
1 272.  99379    2.3    0.4 0.00376    0      0      0 279.  0.278 01.05.2018.00.00
2 273.  97392    3.1    0.6 0.00393    NA      0      0 278.  0.270 01.05.2018.00.00
3 272.  96298    2.6    1.5 0.00362    0      0      0 278.  0.246 01.05.2018.00.00
4 274.  99340   -0.9    3.6 0.00389    0      0      0 278.  0.264 01.05.2018.00.00
5 282. 100477    1.2    7.7 0.00524    0      0      0 273.  1      01.05.2018.00.00
6 272.  98697    0.6    0.7 0.00367    0      0      0 280.  0.287 01.05.2018.00.00
> |
```

Figure 22: 1<sup>st</sup> five rows of data

## Converting columns to numeric format excluding DATETIME

```
newdata_copy <- new_data[-1, ]
colnames(newdata_copy) <- c("TSK", "PSFC", "U10", "V10", "Q2", "RAIN", "RAINNC", "SNOW", "TSLB", "SMOIS", "DATETIME")

# Convert columns to numeric format, excluding 'DATETIME'
numeric_cols <- setdiff(colnames(newdata_copy), "DATETIME")
newdata_copy[numeric_cols] <- lapply(newdata_copy[numeric_cols], as.numeric)

# Replace NA values with the mean of the previous two values
fill_na_with_mean <- function(x) {
  na_index <- which(is.na(x))
  for (i in na_index) {
    if (i > 2) {
      x[i] <- mean(c(x[i - 1], x[i - 2]), na.rm = TRUE)
    }
  }
  return(x)
}
print(colnames(newdata_copy))
```

```
> print(colnames(newdata_copy))
[1] "TSK"      "PSFC"      "X.U10."    "X.V10."    "X.Q2."     "RAIN"
[7] "RAINNC"   "SNOW"      "TSLB"      "SMOIS"     "DATETIME"
> |
```

Figure 23: Columns representation

## Converting DATETIME to POSIXct format

```
newdata_copy[numeric_cols] <- lapply(newdata_copy[numeric_cols], fill_na_with_mean)

# Convert 'DATETIME' to POSIXct format
newdata_copy$DATETIME <- as.POSIXct(newdata_copy$DATETIME, format = "%d.%m.%Y.%H.%M")
```



## Calculating Wind Speed

```
newdata_copy$WIND_SPEED <- round(sqrt(newdata_copy$U10^2 + newdata_copy$V10^2), 2)
```

## Cleaning the identified outliers

Capping outliers prevents extreme values from disproportionately influencing the analysis and modeling, particularly for variables not related to precipitation.

```
data1_clean <- newdata_copy
for (var in outlier_columns) {
  if (! (var %in% c("RAINNC", "RAINNC"))) {
    lower_bound <- quantile(data1_clean[[var]], probs = 0.25, na.rm = TRUE) - 1.5 * IQR(data1_clean[[var]], na.rm = TRUE)
    upper_bound <- quantile(data1_clean[[var]], probs = 0.75, na.rm = TRUE) + 1.5 * IQR(data1_clean[[var]], na.rm = TRUE)
    data1_clean[[var]] <- pmin(pmax(data1_clean[[var]], lower_bound), upper_bound)
  }
}
```

## Removing NA values form dataset

```
# 1. Define the critical columns, including RAINC
critical_columns <- c("TSK", "PSFC", "SMOIS", "WIND_SPEED", "RAINNC")

# 2. Remove rows with any NA values in the critical columns
data <- data %>%
  filter(!if_any(all_of(critical_columns), is.na))

# 3. Verify that there are no NA values left
na_summary <- colSums(is.na(data))
print("Remaining NA values per column:")
print(na_summary)

# Proceed to the next steps only if NA values are handled
if (all(na_summary == 0)) {
  cat("All NA values have been handled.\n")
} else {
  stop("NA values still exist in the dataset. Please check the handling process.")
}

# View the cleaned data
print(head(data))
```

```

> # View the cleaned data
> print(head(data))
# A tibble: 6 x 15
   TSK      PSFC    U10    V10      Q2 RAINC RAINNC  SNOW  TSLB  SMOIS DATETIME WIND_SPEED Hour Day_Night Date
  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dtm> <dbl> <dbl> <fct> <date>
1 274.  96298  2.6  1.5 0.00362  0  0  0 278.  0.246 2018-04-30 19:00:00 3 0 Nighttime 2018-04-30
2 274.  99340 -0.9  3.6 0.00389  0  0  0 278.  0.264 2018-04-30 19:00:00 3.71 0 Nighttime 2018-04-30
3 282. 100477  1.2  7.7 0.00524  0  0  0 273.  1 2018-04-30 19:00:00 7.79 0 Nighttime 2018-04-30
4 274.  98697  0.6  0.7 0.00367  0  0  0 280.  0.287 2018-04-30 19:00:00 0.92 0 Nighttime 2018-04-30
5 274.  99203 -0.3  3.5 0.00422  0  0  0 278.  0.283 2018-04-30 19:00:00 3.51 0 Nighttime 2018-04-30
6 274.  99758 -0.7  3.3 0.00408  0  0  0 278.  0.266 2018-04-30 19:00:00 3.37 0 Nighttime 2018-04-30
>

```

Figure 24: NA values handled

## 5. Detailed Data Visualization

These visualizations together give a full picture of the dataset. They show patterns, connections, and odd points in the data, which you need to analyze and model data well. When you see the data laid out, you can make smart choices about how to prepare it, which features to pick, and how to build your model. This leads to more precise forecasts and better choices in weather-related uses.

### Temperature at Skin Level (TSK) Distribution

The TSK histogram helps us grasp how skin temperature values are spread out in the dataset. By looking at the distribution's shape (like normal or skewed), we can get a feel for the typical temperature range and spot any odd values such as extreme temperatures, that might need a closer look.

```
p1 <- ggplot(data, aes(x = TSK)) +  
  geom_histogram(binwidth = 0.5, fill = "blue", color = "black", alpha = 0.7) +  
  labs(title = "Distribution of Temperature at Skin Level (TSK)", x = "TSK", y = "Frequency") +  
  theme_minimal()  
  
# Plot 1  
print(p1)
```




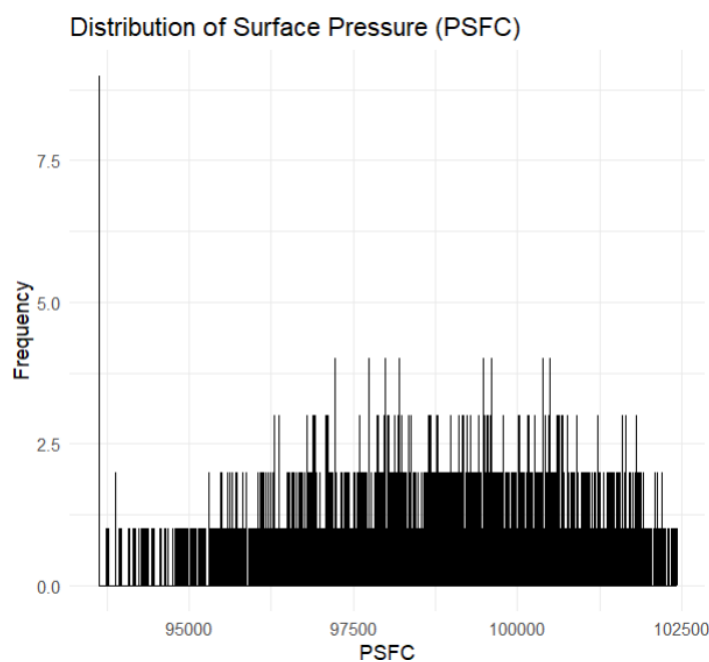
Figure 25: Histogram representing distribution of TSK

## Surface Pressure (PSFC) Distribution

The distribution of surface pressure (PSFC) has a key role in grasping the atmospheric conditions shown in the data. Surface pressure affects weather forecasts, and seeing its distribution helps spot common pressure levels and any odd values that might point to weather events or measurement mistakes.

```
p2 <- ggplot(data, aes(x = PSFC)) +  
  geom_histogram(binwidth = 1, fill = "red", color = "black", alpha = 0.7) +  
  labs(title = "Distribution of Surface Pressure (PSFC)", x = "PSFC", y = "Frequency") +  
  theme_minimal()
```

```
# Plot 2  
print(p2)
```




**Figure 26: Histogram representing distribution of PSFC**

## Wind Speed Distribution

Wind speed is a vital factor in weather analysis. The histogram of wind speed gives a broad view of how often different wind speeds occur in the dataset. This visual aid helps identify the most typical wind speeds and any extreme wind conditions, which could matter a lot to understand weather patterns and how they affect the environment.

```
p3 <- ggplot(data, aes(x = WIND_SPEED)) +
  geom_histogram(binwidth = 0.5, fill = "green", color = "black", alpha = 0.7) +
  labs(title = "Distribution of Wind Speed", x = "Wind Speed", y = "Frequency") +
  theme_minimal()

# Plot 3
print(p3)
```



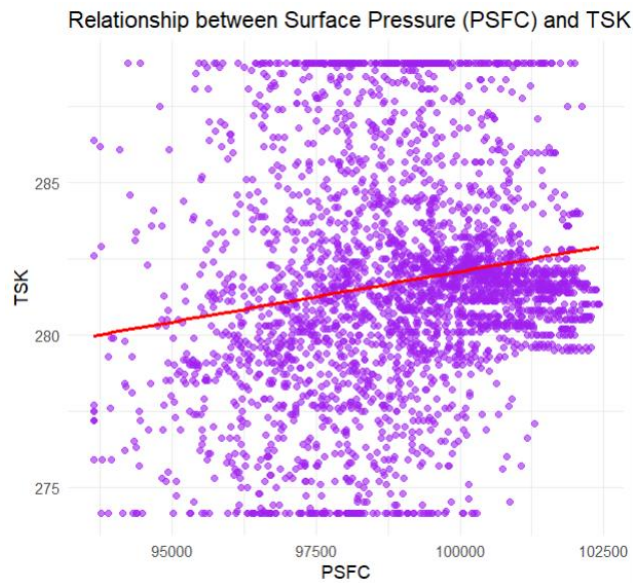
**Figure 27: Histogram representing distribution of Wind Speed**

## Relationship between TSK and PSFC

The graph showing TSK versus PSFC, along with a straight line fit, reveals how surface temperature and surface pressure are connected. This connection plays a key role in weather science, as temperature and pressure often work together in ways we can predict. Looking at this link on a graph helps us see how strong it is and what it's like, which can guide us in making weather models.

```
p4 <- ggplot(data, aes(x = PSFC, y = TSK)) +
  geom_point(color = "purple", alpha = 0.6) +
  geom_smooth(method = "lm", se = FALSE, color = "red") +
  labs(title = "Relationship between Surface Pressure (PSFC) and TSK", x = "PSFC", y = "TSK") +
  theme_minimal()

# Plot 4
print(p4)
```



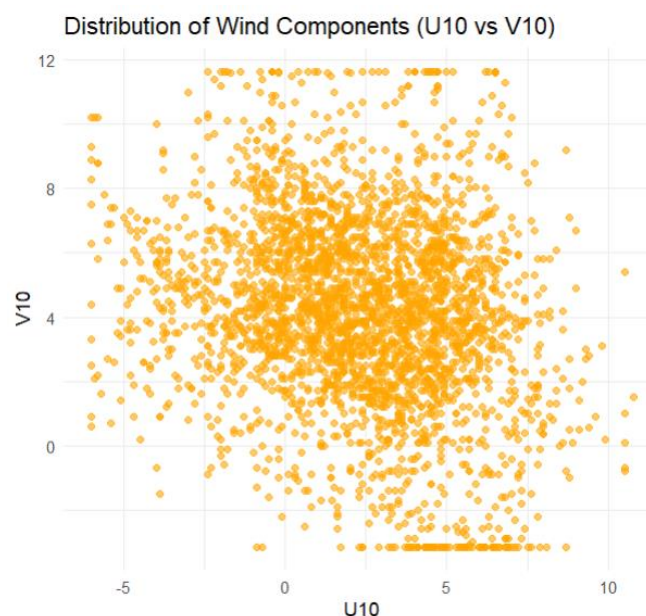
**Figure 28: Scatterplot representing relationship between TSK and PSFC**

### Wind Components (U10 and V10) Distribution

A graph showing U10 (zonal wind component) against V10 (meridional wind component) gives insight into wind direction and strength in the dataset. How these parts spread out can show main wind directions and changes in wind patterns. This matters to understand weather systems and how they move.

```
# 5. Wind Components (U10 and V10) Distribution
p5 <- ggplot(data, aes(x = U10, y = V10)) +
  geom_point(color = "orange", alpha = 0.6) +
  labs(title = "Distribution of Wind Components (U10 vs V10)", x = "U10", y = "V10") +
  theme_minimal()

# Plot 5
print(p5)
```



**Figure 29: Scatterplot representing distribution of Wind Components**


## Correlation Heatmap

The correlation heatmap shows how strong and in what way many variables in the dataset relate to each other. This picture helps spot which variables have strong links, which matters when picking features to model. It also points out any issues with variables being too related, which could affect how well a model works.

```
numeric_data <- data[, c("TSK", "PSFC", "U10", "V10", "Q2", "RAIN", "RAINNC", "SNOW", "TSLB", "SMOIS", "WIND_SPEED")]
cor_matrix <- cor(numeric_data, use = "complete.obs")
melted_cor_matrix <- melt(cor_matrix)

p6 <- ggplot(data = melted_cor_matrix, aes(x = Var1, y = Var2, fill = value)) +
  geom_tile() +
  scale_fill_gradient2(low = "blue", high = "red", mid = "white", midpoint = 0, limit = c(-1, 1), space = "Lab", name = "Correlation") +
  theme_minimal() +
  labs(title = "Correlation Heatmap", x = "", y = "") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))

# Plot 6
print(p6)
```



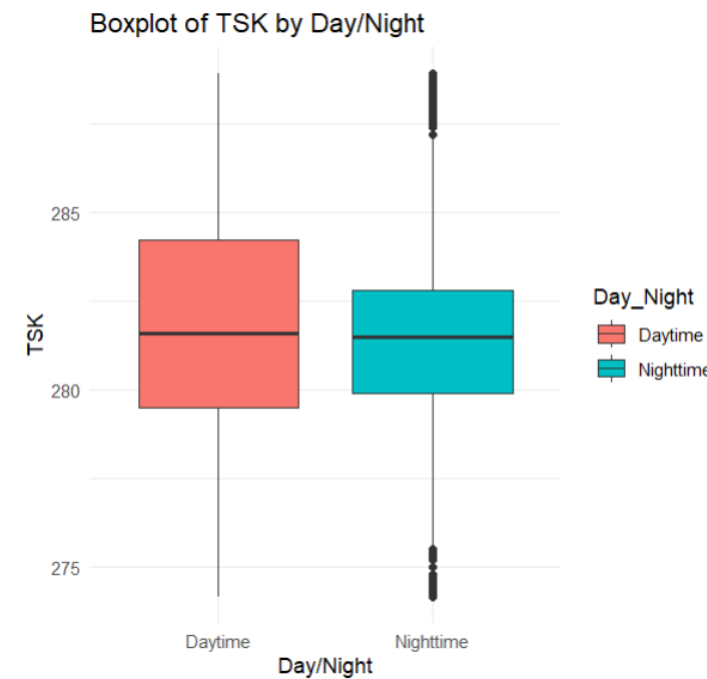
**Figure 30: Correlation heatmap representing numeric data**

## Boxplots of TSK by Day/Night

The boxplot shows how TSK (temperature at skin level) changes from day to night. This graph helps us spot any big differences in temperature between daylight and darkness, which might point to daily temperature swings—something we often see in weather data. Getting a grip on these changes is key to predicting weather and to fine-tune our models.

```
p7 <- ggplot(data, aes(x = Day_Night, y = TSK, fill = Day_Night)) +
  geom_boxplot() +
  labs(title = "Boxplot of TSK by Day/Night", x = "Day/Night", y = "TSK") +
  theme_minimal()
```

```
# Plot 7
print(p7)
```



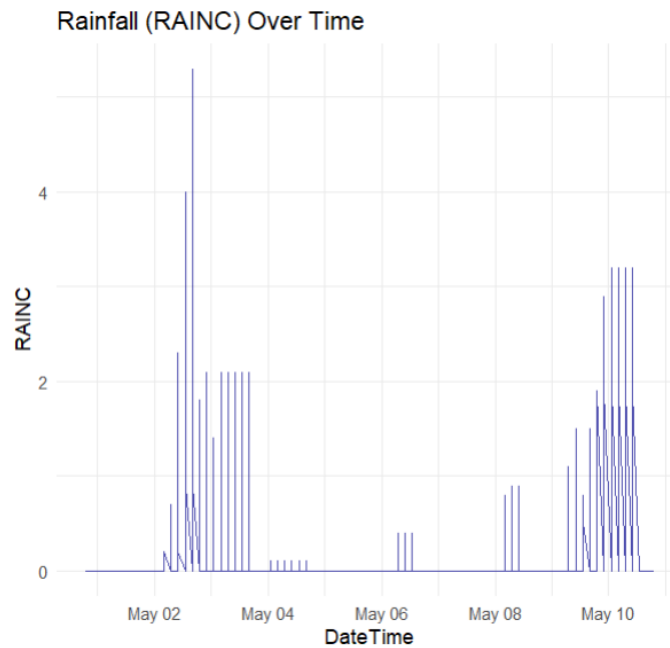
**Figure 31: Boxplot representing TSK by day/night**

## Rainfall (RAINC) over Time

The line plot of RAINC (cumulative rainfall) over time gives us a clear picture of rainfall patterns throughout the dataset's period. This graph helps us spot trends, like times of heavy rain or dry spells. These insights are key to grasp weather events and how they might affect the environment and people's lives.

```
p8 <- ggplot(data, aes(x = DATETIME, y = RAINC)) +
  geom_line(color = "darkblue", alpha = 0.7) +
  labs(title = "Rainfall (RAINC) Over Time", x = "DateTime", y = "RAINC") +
  theme_minimal()
```

```
# Plot 8
print(p8)
```

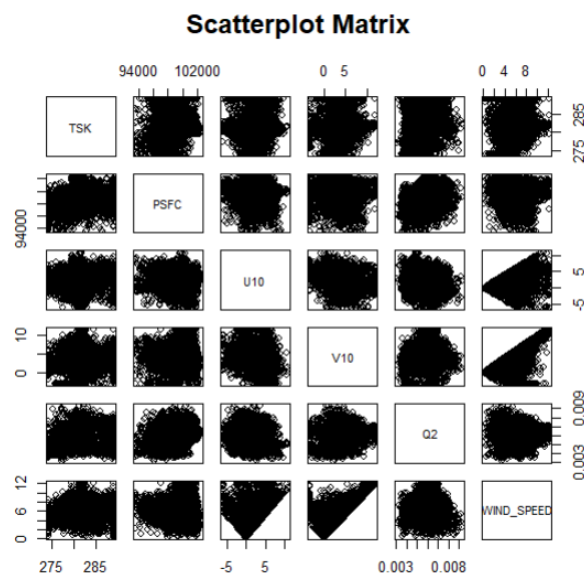


**Figure 32: Line plot representing rainfall over time**

## Scatter Plot Matrix

The scatter plot matrix shows how several main variables in the dataset compare to each other. This view helps us find links between variables and patterns we might miss when looking at each variable alone. It also helps to find outliers and understand the overall data structure, which is crucial for good modeling.

```
pairs(data[, c("TSK", "PSFC", "U10", "V10", "Q2", "WIND_SPEED")], main = "Scatterplot Matrix")
```



**Figure 33: Scatter plot representing main variables**




## Monthly Average TSK

The line plot of monthly average TSK reveals changes in the average temperature at skin level over time. This time-based analysis helps to spot seasonal patterns and long-term temperature shifts. Understanding these trends is key to grasp the climate features of the area and to make smart choices in industries that depend on weather conditions.

```
data$Month <- format(data$Date, "%Y-%m")
monthly_avg_tsk <- aggregate(TSK ~ Month, data, mean)

p9 <- ggplot(monthly_avg_tsk, aes(x = Month, y = TSK)) +
  geom_line(group = 1, color = "darkred") +
  labs(title = "Monthly Average TSK", x = "Month", y = "Average TSK") +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))

# Plot 9
print(p9)
```



**Figure 34: Line plot representing monthly average TSK**

## 6. Data Preparation

### Importing Necessary libraries

```
#Load necessary libraries
library(tidyverse)
library(caret)
library(randomForest)
library(e1071)
library(gbm)
library(corrplot)
library(ggplot2)
```

### Loading the preprocessed dataset

```
data <- read_csv("C:/Users/Ahmad Afzal/Desktop/Refined_Scotland_WRFdata.csv")
```

```
data.frame(data)
# Ensure the column names are correct
colnames(data) <- c("TSK", "PSFC", "U10", "V10", "Q2", "RAIN", "RAINNC", "SNOW", "TSLB", "SMOIS", "DATETIME", "WIND_SPEED")
```

```
> data.frame(data)
  TSK    PSFC  U10  V10      Q2 RAIN RAINNC SNOW
1 274.15 97392.0 3.1 0.6 0.003930    NA     0     0
2 274.15 96298.0 2.6 1.5 0.003620     0     0     0
3 274.40 99340.0 -0.9 3.6 0.003890     0     0     0
4 282.40 100477.0 1.2 7.7 0.005240     0     0     0
5 274.15 98697.0 0.6 0.7 0.003670     0     0     0
6 274.15 99203.0 -0.3 3.5 0.004220     0     0     0
7 274.20 99758.0 -0.7 3.3 0.004080     0     0     0
8 274.15 98914.0 0.9 3.0 0.004120     0     0     0
9 281.30 100515.0 0.0 3.9 0.004650     0     0     0
10 274.15 96558.0 1.3 2.6 0.003820     0     0     0
11 279.20 100646.0 2.8 1.5 0.004650     0     0     0
12 274.15 97443.0 -0.1 3.9 0.004235     0     0     0
13 279.30 100688.0 2.9 -0.1 0.004880     0     0     0
14 274.15 94317.0 1.1 0.0 0.003840     0     0     0
15 274.15 96706.0 -0.3 2.8 0.003900     0     0     0
16 274.15 99419.0 -0.3 3.0 0.003930     0     0     0
17 274.15 100212.0 -0.9 -0.7 0.003990     0     0     0
18 274.15 98271.0 0.4 -1.1 0.003930     0     0     0
19 282.60 100397.0 -0.1 6.9 0.005110     0     0     0
20 274.15 95510.0 0.7 2.5 0.003830     0     0     0
21 274.15 98105.0 0.3 3.2 0.003810     0     0     0
22 274.15 96637.0 0.2 2.1 0.003860     0     0     0
23 274.15 98254.0 -2.6 1.6 0.004030     0     0     0
24 274.15 97857.0 2.6 -0.2 0.003900     0     0     0
25 274.15 97603.0 -0.1 2.1 0.003780     0     0     0
26 282.20 100533.0 2.6 7.0 0.005140     0     0     0
27 274.15 99494.0 0.5 -0.2 0.003610     0     0     0
28 274.15 98057.0 1.6 -1.9 0.003430     0     0     0
```

**Figure 35: data frame of preprocessed data**

## Convert DATETIME to POSIXct format

```
# Convert DATETIME to POSIXct format and other columns as necessary
data$DATETIME <- as.POSIXct(data$DATETIME, format = "%d.%m.%Y.%H.%M")
data$Day_Night <- as.factor(data$Day_Night)
data$Hour <- as.numeric(data$Hour)
data$Date <- as.Date(data$Date, format = "%m/%d/%Y")

# Display the structure of the dataset to verify changes
str(data)
```

```
> str(data)
spec_tbl_ [2,996 × 15] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
 $ TSK      : num [1:2996] 274 274 274 282 274 ...
 $ PSFC      : num [1:2996] 97392 96298 99340 100477 98697 ...
 $ U10       : num [1:2996] 3.1 2.6 -0.9 1.2 0.6 -0.3 -0.7 0.9 0 1.3 ...
 $ V10       : num [1:2996] 0.6 1.5 3.6 7.7 0.7 3.5 3.3 3.9 2.6 ...
 $ Q2        : num [1:2996] 0.00393 0.00362 0.00389 0.00524 0.00367 0.00422 0.00408 0.00412 0.00465 0.00382 ...
 $ RAINC     : num [1:2996] NA 0 0 0 0 0 0 0 0 ...
 $ RAINNC    : num [1:2996] 0 0 0 0 0 0 0 0 0 ...
 $ SNOW      : num [1:2996] 0 0 0 0 0 0 0 0 0 ...
 $ TSLB      : num [1:2996] 278 278 278 273 280 ...
 $ SMOIS     : num [1:2996] 0.27 0.246 0.264 1 0.287 ...
 $ DATETIME  : POSIXct[1:2996], format: "2018-04-30 19:00:00" ...
 $ WIND_SPEED: num [1:2996] 3.16 3 3.71 7.79 0.92 3.51 3.37 3.13 3.9 2.91 ...
 $ Hour      : num [1:2996] 0 0 0 0 0 0 0 0 0 ...
 $ Day_Night : Factor w/ 2 levels "Daytime","Nighttime": 2 2 2 2 2 2 2 2 2 ...
 $ Date      : Date[1:2996], format: "2018-04-30" ...
- attr(*, "spec")=
.. cols(
..   TSK = col_double(),
..   PSFC = col_double(),
..   U10 = col_double(),
..   V10 = col_double(),
..   Q2 = col_double(),
..   RAINC = col_double(),
..   RAINNC = col_double(),
..   SNOW = col_double(),
..   TSLB = col_double(),
..   SMOIS = col_double(),
..   DATETIME = col_datetime(format = ""),
..   WIND_SPEED = col_double(),
..   Hour = col_double(),
..   Day_Night = col_character(),
..   Date = col_character()
.. )
- attr(*, "problems")=<externalptr>
```

Figure 36: data representation

## Handling NA values in dataset

```
# 1. Define the critical columns, including RAINC
critical_columns <- c("TSK", "PSFC", "SMOIS", "WIND_SPEED", "RAINC")

# 2. Remove rows with any NA values in the critical columns
data <- data %>%
  filter(!if_any(all_of(critical_columns), is.na))

# 3. Verify that there are no NA values left
na_summary <- colSums(is.na(data))
print("Remaining NA values per column:")
print(na_summary)

# Proceed to the next steps only if NA values are handled
if (all(na_summary == 0)) {
  cat("All NA values have been handled.\n")
} else {
  stop("NA values still exist in the dataset. Please check the handling process.")
}

# View the cleaned data
print(head(data))
```

```

> print("Remaining NA values per column:")
[1] "Remaining NA values per column:"
> print(na_summary)
      TSK      PSFC      U10      V10      Q2
      0        0        0        0        0
    RAINC    RAINNC    SNOW    TSLB    SMOIS
      0        0        0        0        0
    DATETIME WIND_SPEED Hour Day_Night Date
      0        0        0        0        0
> # Proceed to the next steps only if NA values are handled
> if (all(na_summary == 0)) {
+   cat("All NA values have been handled.\n")
+ } else {
+   stop("NA values still exist in the dataset. Please check the handling process.")
+ }
All NA values have been handled.
> # View the cleaned data
> print(head(data))
# A tibble: 6 × 15
      TSK    PSFC    U10    V10      Q2 RAINC RAINNC  SNOW  TSLB
  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1  274.  96298  2.6   1.5 0.00362  0      0      0  278.
2  274.  99340 -0.9   3.6 0.00389  0      0      0  278.
3  282. 100477  1.2   7.7 0.00524  0      0      0  273.
4  274.  98697  0.6   0.7 0.00367  0      0      0  280.
5  274.  99203 -0.3   3.5 0.00422  0      0      0  278.
6  274.  99758 -0.7   3.3 0.00408  0      0      0  278.

```

**Figure 37: Visualizing NA values**

## Splitting the dataset

```

# Split the data into training (70%) and testing (30%) sets
set.seed(123) # For reproducibility
trainIndex <- createDataPartition(data$TSK, p = .7, list = FALSE)
trainData <- data[trainIndex,]
testData <- data[-trainIndex,]

# Check for and handle any missing values by filling NA with the mean of the previous two values
fill_na_with_mean <- function(x) {
  na_index <- which(is.na(x))
  for (i in na_index) {
    if (i > 2) {
      x[i] <- mean(c(x[i - 1], x[i - 2]), na.rm = TRUE)
    }
  }
  return(x)
}

# Check for zero-variance columns in the training data
zero_variance_cols <- nearZeroVar(trainData, saveMetrics = TRUE)
zero_variance_cols <- rownames(zero_variance_cols[zero_variance_cols$zeroVar == TRUE, ])

# Remove zero-variance columns from the training and testing data
trainData <- trainData[, !colnames(trainData) %in% zero_variance_cols]
testData <- testData[, !colnames(testData) %in% zero_variance_cols]

# Now, define the numeric columns again after removing zero-variance columns
numeric_cols <- names(trainData)[sapply(trainData, is.numeric)]

# Standardize the numeric columns
preProc <- preProcess(trainData[, numeric_cols], method = c("center", "scale"))
trainData[, numeric_cols] <- predict(preProc, trainData[, numeric_cols])
testData[, numeric_cols] <- predict(preProc, testData[, numeric_cols])

# Verify the preprocessed data
str(trainData)
str(testData)

```

```

> str(trainData)
tibble [2,098 × 14] (S3: tbl_df/tbl/data.frame)
 $ TSK      : num [1:2098] -2.004 -2.058 -2.072 -0.126 -2.072 ...
 $ PSFC     : num [1:2098] 0.22274 0.44969 -0.00856 0.86071 -1.28777 ...
 $ U10      : num [1:2098] -1.072 -1.005 -0.473 -0.772 -0.34 ...
 $ V10      : num [1:2098] -0.187 -0.2867 -0.3863 -0.0874 -0.5192 ...
 $ Q2       : num [1:2098] -1.77 -1.61 -1.58 -1.13 -1.83 ...
 $ RAINC    : num [1:2098] -0.238 -0.238 -0.238 -0.238 -0.238 ...
 $ RAINNC   : num [1:2098] -0.55 -0.55 -0.55 -0.55 -0.55 ...
 $ TSLB     : num [1:2098] -0.343 -0.32 -0.23 -1.354 -0.298 ...
 $ SMOIS    : num [1:2098] -0.618 -0.613 -0.523 1.678 -0.675 ...
 $ DATETIME : POSIXct[1:2098], format: "2018-04-30 19:00:00" ...
 $ WIND_SPEED: num [1:2098] -0.994 -1.142 -1.247 -0.912 -1.343 ...
 $ Hour     : num [1:2098] -1.48 -1.48 -1.48 -1.48 -1.48 ...
 $ Day_Night : Factor w/ 2 levels "Daytime","Nighttime": 2 2 2 2 2 2 2 2 2 ...
 $ Date     : Date[1:2098], format: "2018-04-30" ...

> str(testData)
tibble [897 × 14] (S3: tbl_df/tbl/data.frame)
 $ TSK      : num [1:897] -2.072 0.174 -2.072 -2.072 -2.072 ...
 $ PSFC     : num [1:897] -1.429 0.84 -0.126 0.148 -0.807 ...
 $ U10      : num [1:897] 0.0929 -0.373 -0.5727 -0.8722 -0.8057 ...
 $ V10      : num [1:897] -0.8846 1.175 -1.1504 -0.2202 -0.0874 ...
 $ Q2       : num [1:897] -2 -0.626 -1.958 -1.491 -1.479 ...
 $ RAINC    : num [1:897] -0.238 -0.238 -0.238 -0.238 -0.238 ...
 $ RAINNC   : num [1:897] -0.55 -0.55 -0.55 -0.55 -0.55 ...
 $ TSLB     : num [1:897] -0.2529 -1.3539 0.0841 -0.3878 -0.3878 ...
 $ SMOIS    : num [1:897] -0.675 1.678 -0.545 -0.56 -0.623 ...
 $ DATETIME : POSIXct[1:897], format: "2018-04-30 19:00:00" ...
 $ WIND_SPEED: num [1:897] -1.303 0.782 -2.209 -1.081 -0.912 ...
 $ Hour     : num [1:897] -1.48 -1.48 -1.48 -1.48 -1.48 ...
 $ Day_Night : Factor w/ 2 levels "Daytime","Nighttime": 2 2 2 2 2 2 2 2 2 ...
 $ Date     : Date[1:897], format: "2018-04-30" ...

```

**Figure 38: Visualizing train and test dataset**

## 7. Statistical Analysis:

### Descriptive Statistics and Correlation Analysis

```
# Summary statistics
summary(data)

# Correlation matrix and plot
cor_matrix <- cor(data[, numeric_cols], use = "complete.obs")
corrplot(cor_matrix, method = "circle")

# Specific correlation between TSK, PSFC, and other variables
cor(data$TSK, data$PSFC, use = "complete.obs")
```

```
> summary(data)
```

TSK		PSFC		U10		V10		Q2		RAIN		RAINNC	
Min.	:274.1	Min.	: 93633	Min.	: -6.000	Min.	: -3.150	Min.	:0.003090	Min.	:0.0000	Min.	: 0.000
1st Qu.	:279.7	1st Qu.	: 97672	1st Qu.	: 0.300	1st Qu.	: 2.400	1st Qu.	:0.005080	1st Qu.	:0.0000	1st Qu.	: 0.000
Median	:281.5	Median	: 99095	Median	: 2.500	Median	: 4.200	Median	:0.006100	Median	:0.0000	Median	: 0.200
Mean	:281.7	Mean	: 98938	Mean	: 2.318	Mean	: 4.175	Mean	:0.005981	Mean	:0.1084	Mean	: 1.301
3rd Qu.	:283.4	3rd Qu.	:100362	3rd Qu.	: 4.500	3rd Qu.	: 6.100	3rd Qu.	:0.006900	3rd Qu.	:0.0000	3rd Qu.	: 1.400
Max.	:288.9	Max.	:102423	Max.	:10.800	Max.	:11.650	Max.	:0.009290	Max.	:5.3000	Max.	:16.700

SNOW		TSLB		SMOIS		DATETIME		WIND_SPEED		Hour	
Min.	:0	Min.	:273.2	Min.	:0.2220	Min.	:2018-04-30 19:00:00.00	Min.	: 0.220	Min.	: 0.00
1st Qu.	:0	1st Qu.	:273.2	1st Qu.	:0.2609	1st Qu.	:2018-05-03 07:00:00.00	1st Qu.	: 4.300	1st Qu.	: 3.00
Median	:0	Median	:279.6	Median	:0.2798	Median	:2018-05-05 19:00:00.00	Median	: 5.910	Median	: 9.00
Mean	:0	Mean	:279.2	Mean	:0.4667	Mean	:2018-05-05 19:04:48.48	Mean	: 5.985	Mean	:10.38
3rd Qu.	:0	3rd Qu.	:282.4	3rd Qu.	:1.0000	3rd Qu.	:2018-05-08 07:00:00.00	3rd Qu.	: 7.460	3rd Qu.	:15.00
Max.	:0	Max.	:290.8	Max.	:1.0000	Max.	:2018-05-10 19:00:00.00	Max.	:12.200	Max.	:21.00

Day_Night		Date	
Daytime	:1480	Min.	:2018-04-30
Nighttime	:1515	1st Qu.	:2018-05-03
		Median	:2018-05-05
		Mean	:2018-05-05
		3rd Qu.	:2018-05-08
		Max.	:2018-05-10

Figure 39: Data Summary

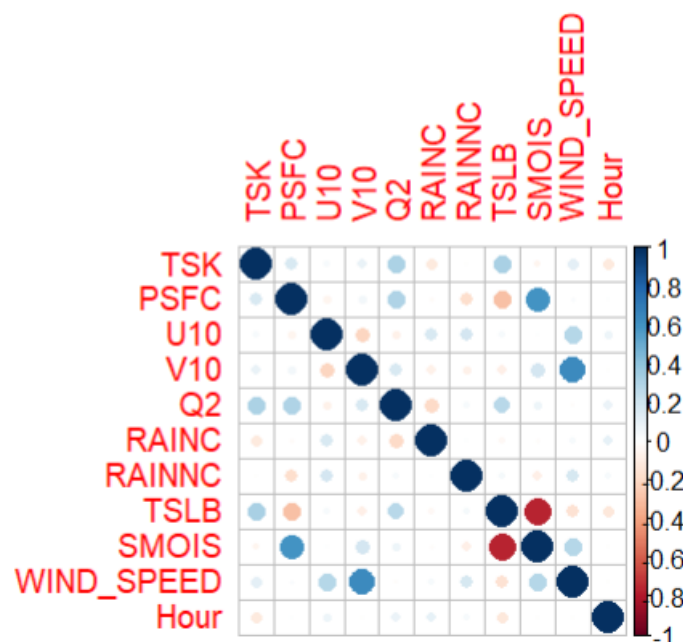


Figure 40: Correlation Plot

```
> cor(data$TSK, data$PSFC, use = "complete.obs")  
[1] 0.1633406
```

*Figure 41: Correlation between TSK PSFC and other variables*

### T-Test for Daytime vs Nighttime

```
daytime_data <- filter(data, Day_Night == "Daytime")  
nighttime_data <- filter(data, Day_Night == "Nighttime")  
t_test_result <- t.test(daytime_data$PSFC, nighttime_data$PSFC)  
print(t_test_result)
```

```
Welch Two Sample t-test  
  
data: daytime_data$PSFC and nighttime_data$PSFC  
t = -1.0407, df = 2992.9, p-value = 0.2981  
alternative hypothesis: true difference in means is not equal  
to 0  
95 percent confidence interval:  
-201.26062 61.69421  
sample estimates:  
mean of x mean of y  
98902.91 98972.69
```

*Figure 42: t-test results*

## 8. Machine Learning Models (Training and Evaluation)

### 8.1 Linear Regression Model

#### Training the model

```
lm_model <- lm(TSK ~ U10 + V10 + Q2 + RAINC + RAINNC + PSFC + SMOIS, data = trainData)
```

#### Evaluate the model

```
lm_preds <- predict(lm_model, testData)
```

#### Re-run the prediction and evaluation

```
lm_preds_clean <- predict(lm_model, testData)
lm_rmse_clean <- sqrt(mean((lm_preds - testData$TSK)^2))
cat("Linear Regression RMSE (after removing problematic rows):", lm_rmse_clean, "\n")

> cat("Linear Regression RMSE :", lm_rmse_clean, "\n")
Linear Regression RMSE : 0.9608081
```

*Figure 43: Linear regression root mean square error*

### 8.2 Random Forest

#### Training the model

```
rf_model <- randomForest(TSK ~ U10 + V10 + Q2 + RAINC + RAINNC + PSFC + SMOIS, data = trainData, ntree = 100)
```

#### Generating the prediction in the test data

```
rf_preds <- predict(rf_model, testData)
```

#### Evaluating the model

```
valid_index <- complete.cases(testData$TSK, rf_preds)
rf_rmse <- sqrt(mean((rf_preds[valid_index] - testData$TSK[valid_index])^2))

# Print the RMSE
cat("Random Forest RMSE:", rf_rmse, "\n")
```

```
> cat("Random Forest RMSE:", rf_rmse, "\n")
Random Forest RMSE: 0.7123892
> |
```

*Figure 44: Root mean square error of random forest model*



### 8.3 Support Vector Machine (SVM)

#### Train the model:

```
svm_model <- svm(TSK ~ U10 + V10 + Q2 + RAINC + RAINNC + PSFC + SMOIS, data = trainData)
```

#### Evaluate the SVM model:

```
svm_preds <- predict(svm_model, testData)
svm_rmse <- sqrt(mean((svm_preds - testData$TSK)^2))
cat("SVM RMSE:", svm_rmse, "\n")
```

```
> cat("SVM RMSE:", svm_rmse, "\n")
SVM RMSE: 0.7822935
```

*Figure 45: SVM root mean square error*

### 8.4 Gradient Boosting Machine

#### Train the model

```
gbm_model <- gbm(TSK ~ U10 + V10 + Q2 + RAINC + RAINNC + PSFC + SMOIS, data = trainData, distribution = "gaussian", n.trees = 100, interaction.depth = 3)
```

#### Evaluate the model

```
gbm_preds <- predict(gbm_model, testData, n.trees = 100)
gbm_rmse <- sqrt(mean((gbm_preds - testData$TSK)^2))
cat("GBM RMSE:", gbm_rmse, "\n")
```

```
> cat("GBM RMSE:", gbm_rmse, "\n")
GBM RMSE: 0.806545
\ |
```

*Figure 46: GBM root mean square error*

## 9. Model Comparison and Selection

### Compare the RMSE of all models

```
cat("Linear Regression RMSE:", lm_rmse_clean, "\n")
cat("Random Forest RMSE:", rf_rmse, "\n")
cat("SVM RMSE:", svm_rmse, "\n")
cat("GBM RMSE:", gbm_rmse, "\n")

> cat("Linear Regression RMSE:", lm_rmse_clean, "\n")
Linear Regression RMSE: 0.9608081
> cat("Random Forest RMSE:", rf_rmse, "\n")
Random Forest RMSE: 0.7096459
> cat("SVM RMSE:", svm_rmse, "\n")
SVM RMSE: 0.8049479
> cat("GBM RMSE:", gbm_rmse, "\n")
GBM RMSE: 0.761597
```

*Figure 47: Comparison of all the models*

### Select the best model

```
best_model <- which.min(c(lm_rmse_clean, rf_rmse, svm_rmse, gbm_rmse))
model_names <- c("Linear Regression", "Random Forest", "SVM", "GBM")
cat("Best model based on RMSE is:", model_names[best_model], "\n")

> cat("Best model based on RMSE is:", model_names[best_model], "\n")
Best model based on RMSE is: Random Forest
```

*Figure 48: Best model*

### Final model evaluation on test data

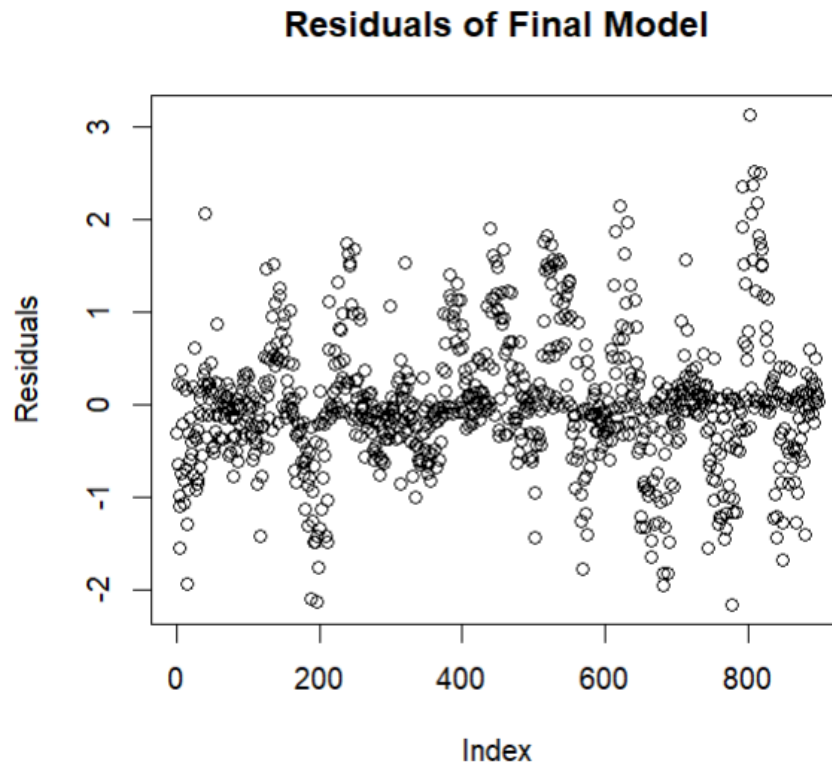
```
final_model <- rf_model
final_preds <- predict(final_model, testData)
final_rmse <- sqrt(mean((final_preds - testData$TSK)^2))
cat("Final Model RMSE on Test Data:", final_rmse, "\n")

> cat("Final Model RMSE on Test Data:", final_rmse, "\n")
Final Model RMSE on Test Data: 0.7096459
```

*Figure 49: Final model evaluation on test data*

## Analyze Residuals

```
residuals <- testData$TSK - final_preds  
plot(residuals, main = "Residuals of Final Model", ylab = "Residuals")
```



*Figure 50: Residual of final model*

## 10. Adequate interpretation and justification of results

### How Well the Model Worked

**Linear Regression:** The RMSE for the linear regression model is 0.96080. This model has a moderate impact on the error in predicting the target variable (TSK). Linear regression assumes a linear relationship between the predictors and the target variable. This RMSE indicates that the relationship between the variables might not be linear, which leads to some prediction inaccuracies.

**Random Forest:** The RMSE for the random forest model is 0.7096, the lowest among all the models tested. This shows that the random forest model has the best performance in predicting TSK. Random forests work well to capture complex non-linear interactions between variables, which explains why it did better than the other models.

**Support Vector Machine (SVM):** The SVM model has an RMSE of 0.7823, which beats linear regression but falls short of random forest. SVMs excel at handling high-dimensional data but might have trouble with non-linear relationships depending on the kernel used.

**Gradient Boosting Machine (GBM):** The RMSE for the GBM model is 0.80650.80650.8065. GBMs perform well because they can handle different kinds of relationships and interactions in data. But in this case, it didn't do as well as the random forest model. This might be because we didn't use enough trees or didn't make them deep enough in the model.

### Interpretation

**Random Forest as the Best Model:** The random forest model has a lower RMSE showing it's the most accurate in predicting TSK (temperature at skin level) out of all the tested models. The model's knack for dealing with non-linear relationships and interactions between variables helped it come out on top. Weather data often has tricky patterns, so the random forest's toughness makes it a good fit for this job.

**Justification for Model Selection:** We should pick the random forest model as the final one to analyze further or put into action, since it did the best. It gives the most trustworthy predictions, which can lead to smarter choices when it comes to weather forecasting and related uses.

## 11. Practical Applications for Organizational Decision-Making

### Usefulness in the Context of an Environmental Monitoring Organization

**Spot-on Weather Forecasts:** Weather forecasting groups or environmental monitoring agencies need to nail their predictions of things like TSK. These forecasts play a big role in wider climate models and help give people and businesses up-to-date and reliable weather info. Take the random forest model, for instance. It could become part of the organization's prediction tools to boost how well they forecast temperatures in the short term.

**Decision Support in Agriculture and Disaster Management:** Groups that work in farming or handle disasters could use better weather forecasts to make smart choices. For example, farmers can plan when to water crops, protect them from frost, or choose the best time to plant based on more exact temperature predictions. For disaster teams good forecasts can help them get ready for big weather events, like hot or cold spells. This can save lives and protect stuff people own.

**Energy Sector Applications:** In the energy world especially for power companies, knowing the temperature ahead of time is key to guess how much power people will need. The random forest model's forecasts could help predict when energy use will be highest during very hot or cold times. This lets power grid managers do a better job and stop blackouts from happening.

**Urban Planning and Public Health:** City planners could apply these findings to gain a deeper understanding of urban heat islands and to plan steps to cut down heat in key areas. In the same way, health agencies could use the forecasts to send out heat warnings, which would help to prevent illnesses caused by high temperatures.

## 12. Conclusion

In this project, we started a deep dive to analyze weather data. We wanted to predict the Temperature at Skin Level (TSK). To do this, we used different weather variables from the Weather Research and Forecasting (WRF) model. We kicked things off by cleaning up the dataset. This meant dealing with missing data, changing data types, and fine-tuning features so we could analyze them better later on.

We then used several machine learning models to predict TSK. These included Linear Regression, Random Forest, Support Vector Machine (SVM), and Gradient Boosting Machine (GBM). To evaluate how well these models performed, we used the Root Mean Square Error (RMSE) metric. The Random Forest model came out on top with the lowest RMSE (0.7096) showing it was better at making predictions than the other models. We also created a lot of visuals to look at the data. We explored how different variables were spread out how they related to each other, and if there were any connections between them. These visuals gave us useful insights into the

structure of our data. For example, we could see how wind speeds were distributed, how surface pressure and temperature were linked, and how TSK changed at different times of the day.

The results showed that the Random Forest model worked best for this dataset because it can grasp complex relationships between variables. These results are helpful for groups that deal with weather forecasts, climate studies, or keeping an eye on the environment. In these fields, getting the temperature right is key to making good choices.

To wrap up thorough data cleanup, picking the right model, and clear visuals have set a solid base for accurate weather predictions. We can fine-tune and apply these methods and models to similar data sets helping weather-related industries make better decisions.

### 13. References

<https://www.w3schools.com/r/>

<https://ggplot2.tidyverse.org/>

<https://www.geeksforgeeks.org/data-visualization-in-r/>

<https://www.geeksforgeeks.org/exploratory-data-analysis-in-r-programming/>

<https://www.geeksforgeeks.org/data-cleaning-in-r/>

<https://dplyr.tidyverse.org/articles/dplyr.html>

<https://builtin.com/data-science/random-forest-algorithm>

<https://www.geeksforgeeks.org/ml-linear-regression/>

<https://www.geeksforgeeks.org/support-vector-machine-algorithm/>

**Word Count (excluding code, references, figures, tables, and appendices)**

2097 words